



Your AI should be a staff, not a search box.

Many purpose-built agents · on hardware you own · thinking with the AI you choose · reachable by the people you let in · through Telegram · on one subscription.

$$\text{OpenClaw} + \text{Docker} + \text{Telegram} = \text{Hatchabot}$$

the open-source agent runtime · one isolated container per agent · the front door everyone already has

THE PROBLEM

A brilliant model, in the wrong shape

You already pay for a model that can do almost anything. What you get is **one assistant, in one app, from one company, for one person** — whose memory of you is whatever that company decides to keep.

- **“I keep re-explaining myself.”** Every conversation starts from zero.
- **“One assistant can't be good at everything.”** Your tax questions, your kids' homework and your condo board all get the same generic voice.
- **“I don't want to be locked in.”** Your history, your prompts and your habits live in someone else's product.
- **“Only I can use it.”** Your family can't message it, and it does nothing when you're not looking.

Chat apps are sessions you visit

Open the app, type, close it, forget. The model is powerful; the container around it is a search box with a scroll bar.

Hatchabot agents are things you own

Files on your disk with a job, a personality and a memory that grows for years — behind a contact your whole household can message.

THE IDEA

The opposite shape

Many agents, each with a job

A kitchen helper, a tax advisor, a homework tutor, a scheduler that reads its own inbox. Made in a minute: a name, a paragraph, a bot token.

On a computer you own

Every agent is a container on your machine. Memory, credentials and history stay home. Pair it with a local model and nothing leaves at all.

With the AI you choose

Claude (subscription or API), Gemini or a local model — per agent, switchable live. Cheap models for simple jobs, the best for the ones that matter.

For the people you let in

Invite family by link or QR. Talk to one agent privately or together in a group room, with memory shared or personal — and everyone told which.

Through Telegram

The app everyone already has. Each agent is a real contact with a name and a face. Nothing to install, no accounts to create.

Run as a fleet

Not a chatbot but a control plane: health, usage, audit trail, backups, security posture, upgrades rolled out one agent at a time.

OpenClaw + Docker = Hatchabot. Every agent is a stock OpenClaw gateway — the open-source agent runtime — in its own Docker container with its own volume, bot and credentials. Hatchabot is OpenClaw *new and improved for a household*: many isolated installs, run as a fleet, by more than one person.

What we believe an agent is

Agents are things you own

Each one is a container with a persona, operating rules and a memory — plain files you can read and edit, snapshotted before every change:

`SOUL.md` · `AGENTS.md` · `MEMORY.md`

An agent survives restarts, rebuilds, model switches and moves between machines.

Memory is a file, not a chat log

The agent writes durable facts as it goes. Save a conversation to memory, download its whole chat history, or recover the context of a thread that got reset.

Nothing important is trapped in a transcript.

Not locked to one lab

Anthropic, Google, or a local server with no account at all — per agent, switchable live without a rebuild.

Group agents into **classes** and retune a whole tier in one place. Your agents outlive any provider's pricing decision.

An agent is a folder of files you own — not a session you visit. What it learns about you accumulates for years, and it's yours.

How it should behave

The front door is Telegram

No app to ship, no accounts to run, no notification system to build. Bots are first-class in Telegram: deep links, group chats and per-agent allowlists, for free.

Isolation by construction

One container, one bot, one set of credentials per agent. Give an agent exactly the data it needs — a folder, a git repo, a Google account with “read mail but never send” — and nothing else.

A daily posture check names every agent that has both a wide audience and a powerful capability.

Reversible, or honest about it

Snapshots before edits and rebuilds. Rollback on a failed move or deploy. A full audit timeline. Usage reported as tokens, not a made-up bill. Health probes that tell “listed as running” from “actually answering”.

One agent failing — a runaway task, a bad edit, a rate-limited source — never takes the others with it.

TWO DELIBERATE CHOICES

Claude Max and Telegram, on purpose

1 plan

One subscription, the whole household

A Claude Max plan is a flat monthly fee. `claude setup-token` turns it into a credential Hatchabot keeps on your machine and hands to every agent — revocable, never a shared login.

The kids' tutor, the family planner, your own advisors: **everyone gets their own agents, and you pay for one plan.** Nobody else needs an AI account — they message the agents made for them.

20 bots

Telegram: 20 real identities per free account

Free, on every platform, and bots are a first-class feature. BotFather makes one in a minute; a free account can own 20, Premium 40, and every household member adds more.

The alternative — our own app, our own accounts, our own push notifications — would cost more than the whole rest of the product.

One plan and a free front door: a household of agents that's affordable to leave on.

FEATURES

What you can do that you can't do with a chat app

CAPABILITY	IN PRACTICE
Create agents fast	A name, one paragraph and a bot token → a running, remembering agent in a minute. Clone one you like; import a shared template; keep a launchpad of planned agents.
Many humans, one agent	Private one-to-one threads with the same agent, or a shared room. Invite by link or QR. Shared or private memory, declared.
Choose the brain per agent	Claude, Gemini or a local model; switch the model live; classes for tiers; migrate a whole source at once.
Give agents accounts and data	Gmail, Calendar, Drive and Sheets per agent (with “no send”); read-only or writable folders; git repos it commits to — public repos need no key at all; per-agent secrets.
Schedule and react	Cron tasks, event-triggered tasks (a zero-cost check decides whether to wake the model), run-now tests.
Agents that consult each other	Grant an investing agent access to your tax and legal agents; it asks them mid-task, with scoped tokens, loop guards and rate limits.
Master → child lineage	A master agent's improvements distil into proposals that push to its children.
Move, copy, share, back up	Rehost to another machine in one step; download as one file; share as a template with no secrets; nightly backups with a restore drill.
Rebuild without fear	Rebuilds keep memory. Only a change of AI source resets a thread — and that offers to save memory first, and to recover the context after.

Operator-grade, because agents that matter need it

See what's going on

- Fleet health: probed, not assumed
- Usage per agent per day, tokens per hour
- An audit timeline of every change
- A daily security posture with diffs
- A legend showing every pending action

Change things safely

- Bulk actions by source, group, class or state
- Classes: model + source + runtime image, retuned in one place
- Runtime images: build a candidate, try it on **one** agent, promote or discard
- Snapshots before every edit; rollback on failed moves

Keep it yours

- Nightly backups: every volume, the registry, the key
- Restore one agent from any backup set
- Export and import as a single file
- Runners: place agents on a second machine
- All over a private Tailscale network — nothing opened to the internet

Everything an operator needs — and nothing a family member ever has to see.

COMPARED

Versus Claude.ai, ChatGPT and Codex

A chat app is a place you go to talk to one assistant. Hatchabot is a staff of agents that live on your machine and come to you.

	CLAUDE.AI / CHATGPT	CODEX (AND CODING AGENTS)	HATCHABOT
Shape	One assistant, one app, one person	One agent, one repo, one task at a time	A fleet of agents with jobs, for a household
Runs where	Their cloud	Their cloud or your terminal	Your machine (containers), over your private network
Memory	What the provider keeps; opaque	The repo and the session	Files you read, edit, back up; survives rebuilds and moves
Who can talk to it	You, in the app	You, in the terminal or a PR	Anyone you invite, on Telegram; private threads or group rooms
Choice of AI	That company's models	That company's models	Claude, Gemini, local — per agent, switchable live
Acts on its own	Limited scheduled tasks	When you run it	Schedules, events, and requests from other agents
Access to your things	Connectors the vendor chose	The repo	Per agent: a folder, a git repo, a Google account — limits you set
Cost	Per seat	Per token / per seat	One Claude Max plan powers every agent; local models cost nothing per token
Lock-in	Your history stays with them	The repo is yours; the service is theirs	Export, move machines, change provider — the agents are yours

The honest caveat: the chat apps are polished, need no setup, and their models are excellent. Hatchabot doesn't compete with the model — it *uses* one (Claude by default). It competes with the app around the model: one assistant, one person, forgets you tomorrow. A coding agent with a git repo is one of the things you can make.

COMPARED

OpenClaw, new and improved: what Hatchabot adds

Every Hatchabot agent **is** an OpenClaw gateway. Hatchabot adds the layer OpenClaw deliberately leaves out: many installs, run as a fleet, by more than one person.

One install per agent, so agents can't hurt each other

Plain OpenClaw is one process, one config, one workspace and credential set shared by every agent in it. Hatchabot gives each agent its own container, volume, bot token, tool policy and secrets.

Upgrades the same way

The runtime image is pinned and rolled out per agent, candidate first — never “upgrade the gateway and hope every agent survives”.

Lifecycle instead of hand-editing

Create, clone, rebuild, archive, restore, delete, move — with a snapshot before every change. In plain OpenClaw that's you, a shell and a directory.

More than one human

Owners, members, invites, pairing approvals, group rooms, and memory people are told is shared or private.


Credentials managed once

One setup-token serves every agent; Google accounts attach per agent; secrets are write-only.

If you run one agent for yourself, plain OpenClaw is fine. Hatchabot is for when there are several, they matter, and other people talk to them.

HOW IT WORKS

One control plane, one container per agent



Persisted

Each agent's `$HOME` is a Docker volume: persona, memory, sessions, tools. The runtime image is shared, versioned, and swapped under an agent without touching the volume.

Private by default

The web app lives on your machine; reach it from anywhere over Tailscale, never a port opened on your router. Telegram works from anywhere regardless.

Open source

TypeScript, MIT licensed, 965 automated tests and fifteen full audits. Tagged releases with notes; pre-built images for arm64 and amd64.

BUILT TO BE EASY TO RUN

From a spare computer to a running agent

1 line

One-line installer

Checks git, Docker and Node, offers to install what's missing, fetches the latest release and installs the background service.

~1 min

Pre-built runtime image

Published for arm64 and amd64 with every release, so setup pulls the agent image instead of building it for twenty minutes.

3 steps

Setup wizard

The app walks you through connecting the AI, making the first Telegram bot and creating the first agent — each step checked as you go.

`hatchabot doctor`

Checks Node, Docker, the image, the service, the database, disk, backups and Tailscale — and prints the fix for anything wrong.

Upgrades by tag, with rollback

Every release is a tag with notes. Deploying one restarts only the control plane; if it doesn't come up healthy, the previous release is put back.

Backups on by default

Nightly sets of every agent, the registry and the key, kept for two weeks, with a restore drill that proves a set actually restores.

GET STARTED

Your first agent in 15 minutes

- 1 Make a Telegram bot.** In Telegram, open @BotFather, send `/newbot`, and keep the token.
- 2 Get Claude ready.** Install the Claude CLI, log in once, then run `claude setup-token` and keep that token too.
- 3 Install Hatchabot** on the computer that stays on:

```
$ bash -c "$(curl -fsSL https://hatchabot.com/install.sh)"
```
- 4 Open the app** at `http://localhost:8080`. The wizard takes both tokens and creates your first agent.
- 5 Say hi** on Telegram. It will still remember this conversation next year.

You need

- A Linux or macOS computer that stays on — a Mac mini, a home server, an old laptop
- A Telegram account
- A Claude subscription (Pro for a couple of agents, Max for a household) — or an API key, or a local model

Then, when you like

- Tailscale, to open the app from anywhere
- Invite family by link or QR
- Google accounts, folders, repos, schedules, peers



**For people who want more than one assistant
— shared with a household, on their own
machine, with the AI they choose.**

```
$ bash -c "$(curl -fsSL https://hatchabot.com/install.sh)"
```

hatchabot.com · github.com/hatchabot/hatchabot · hello@hatchabot.com · MIT licensed